

get_homologues-est manual

Bruno Contreras-Moreira (1,2) and Pablo Vinuesa (3)

1. Fundación ARAID and 2. Estación Experimental de Aula Dei-CSIC
3. Centro de Ciencias Genómicas, Universidad Nacional Autónoma de México

November 13, 2017

Contents

1	Description	3
2	Requirements and installation	3
2.1	Perl modules	4
2.2	Required binaries	4
2.3	Optional software dependencies	5
3	User manual	8
3.1	Input data	8
3.2	Program options	9
3.3	Accompanying scripts	14
4	A few examples	16
4.1	Extracting coding sequences (CDS) from transcripts	16
4.2	Clustering orthologous transcripts from FASTA files, one per strain	18
4.3	Producing a nucleotide-based pangenome matrix	23
4.4	Estimating protein domain enrichment of some sequence clusters	24
4.5	Making and annotating a non-redundant pangenome matrix	26
4.6	Annotating a sequence cluster	28
5	A step-by-step protocol with barley assembled transcripts	30
6	Frequently asked questions (FAQs)	33
7	Credits and references	37

1 Description

This document describes *get_homologues-est*, a fork of *get_homologues* for clustering homologous gene/transcript sequences of strains/populations of the same species. This algorithm has been designed and tested with plant transcripts and CDS sequences, and uses BLASTN to compare DNA sequences. The main tasks for which this was conceived are:

- Finding and translating coding regions (CDSs) within raw transcripts.
- Clustering transcripts/CDS nucleotide sequences in homologous (possibly orthologous) groups, on the grounds of DNA sequence similarity.
- Definition of pan- and core-transcriptomes by calculation of overlapping sets of CDSs.

The core algorithms of *get_homologues-est* have been adapted from *get_homologues*, and are therefore explained in [manual_get_homologues.pdf](#). This document focuses mostly on EST-specific options.

This table lists features developed for *get_homologues-est* which were not available in the original *get_homologues* release, although most have been backported.

name	description
transcripts2cds.pl	Script to extract coding sequences CDS from raw transcripts by combining Transdecoder and BLASTX.
redundant isoform calling	<i>get_homologues-est</i> can handle redundant isoforms which otherwise will degrade clustering performance.
ANI matrices	<i>get_homologues-est</i> can compute Average Nucleotide Identity (ANI) matrices which summarize the genetic distance among input genotypes.
make_nr_pangenome_matrix.pl	Produces a non-redundant pangenome matrix by comparing all nucleotide/peptide clusters to each other.
pfam_enrich.pl	Script to test whether a set of sequence clusters are enriched in some Pfam domains.
annotate_cluster.pl	Script to reconstruct the BLAST alignments that support a selected sequence cluster.

Table 1: List of novel scripts/features in *get_homologues-est*.

2 Requirements and installation

get_homologues-est.pl is a Perl5 program bundled with a few binary files. The software has been tested on 64-bit Linux boxes, and on Intel MacOSX 10.11.1 systems. Therefore, a Perl5 interpreter is needed to run this software, which is usually installed by default on these operating systems.

In order to install and test this software please follow these steps:

1. Unpack the software with: `$ tar xvfz get_homologues_X.Y.tgz`
2. `$ cd get_homologues_X.Y`
3. `$./install.pl`
Please follow the indications in case some required part is missing.
4. Type `$./get_homologues-est.pl -v` which will tell exactly which features are available.
5. Test the main Perl script, named `get_homologues-est.pl`, with the included sample input folder `sample_transcripts_fasta` by means of the instruction:
`$./get_homologues-est.pl -d sample_transcripts_fasta`. You should get an output similar to the contents of file `sample_transcripts_output.txt`.
6. Optionally modify your `$PATH` environment variable to include *get_homologues-est.pl*. Please copy the following lines to the `.bash_profile` or `.bashrc` files, found in your home directory, replacing `[INSTALL_PATH]` by the full path of the installation folder:

```
export GETHOMS=[INSTALL_PATH]/get_homologues_X.Y
export PATH=${GETHOMS}/:${PATH}
```

This change will be effective in a new terminal or after running: `$ source ~/.bash_profile`

The rest of this section might be safely skipped if installation went fine, it was written to help solve installation problems.

2.1 Perl modules

A few Perl core modules are required by the *get_homologues-est.pl* script, which should be already installed on your system: Cwd, FindBin, File::Basename, File::Spec, File::Temp, FileHandle, List::Util, Getopt::Std, Benchmark and Storable.

In addition, the *Bio::Seq*, *Bio::SeqIO*, *Bio::Graphics* and *Bio::SeqFeature::Generic* modules from the [Bioperl](#) collection, and module [Parallel::ForkManager](#) are also required, and have been included in the *get_homologues-est* bundle for your convenience.

Should this version of BioPerl fail in your system (as diagnosed by *install.pl*) it might be necessary to install it from scratch. However, before trying to download it, you might want to check whether it is already living on your system, by typing on the terminal:

```
$ perl -MBio::Root::Version -e 'print $Bio::Root::Version::VERSION'
```

If you get a message `Can't locate Bio/Root/Version...` then you need to actually install it, which can sometimes become troublesome due to failed dependencies. For this reason usually the easiest way of installing it, provided that you have root privileges, it is to use the software manager of your Linux distribution (such as *synaptic/apt-get* in Ubuntu, *yum* in Fedora or *YaST* in openSUSE). If you prefer the terminal please use the *cpan* program with administrator privileges (*sudo* in Ubuntu):

```
$ cpan -i C/CJ/CJFIELDS/BioPerl-1.6.1.tar.gz
```

This form should be also valid:

```
$ perl -MCPAN -e 'install C/CJ/CJFIELDS/BioPerl-1.6.1.tar.gz'
```

Please check this [tutorial](#) if you need further help.

2.2 Required binaries

In order to properly read (optionally) compressed input files, *get_homologues-est* requires *gunzip* and *bunzip2*, which should be universally installed on most systems.

The Perl script *install.pl*, already mentioned in section 2, checks whether the included precompiled binaries for [hmmer](#), [MCL](#) and [BLAST](#) are in place and ready to be used by *get_homologues-est*. This includes also [COGtriangles](#), which is used only by prokaryotic *get_homologues*. However, if any of these binaries fails to work in your system, perhaps due a different architecture or due to missing libraries, it will be necessary to obtain an appropriate version for your system or to compile them with your own compiler.

In order to compile *MCL* the GNU *gcc* compiler is required, although it should most certainly already be installed on your system. If not, you might install it by any of the alternatives listed in section 2.1. For instance, in Ubuntu this works well: `$ sudo apt-get install gcc`. The compilation steps are as follows:

```
$ cd bin/mcl-14-137;
$ ./configure';
$ make
```

To compile *COGtriangles* the GNU *g++* compiler is required. You should obtain it by any of the alternatives listed in section 2.1. The compilation would then include several steps:

```
$cd bin/COGsoft;
$cd COGlse; make;
$cd ../COGmakehash;make;
$cd ../COGreadblast;make;
$cd ../COGtriangles;make
```

Regarding BLAST, *get_homologues-est* uses BLAST+ binaries, which can be easily downloaded from the [NCBI FTP](#) site. The packed binaries are *blastp* and *makeblastdb* from version *ncbi-blast-2.2.27+*. If these do not work

in your machine or you prefer to use older BLAST versions, then it will be necessary to edit file *lib/phyTools.pm*. First, environmental variable `$ENV{'BLAST_PATH'}` needs to be set to the right path in your system (inside subroutine `sub set_phyTools_env`).

Variables `$ENV{'EXE_BLASTP'}` and `$ENV{'EXE_FORMATDB'}` also need to be changed to the appropriate BLAST binaries, which are respectively *blastall* and *formatdb*.

2.3 Optional software dependencies

It is possible to make use of *get_homologues-est* on a computer farm or high-performance computing cluster managed by [gridengine](#). In particular we have tested this feature with versions GE 6.0u8, 6.2u4, 2011.11p1 invoking the program with option `-m cluster`. For this command to work it might be necessary to edit the *get_homologues-est.pl* file and add the right path to set global variable `$SGEPATH`. To find out the installation path of your SGE installation you might try the next terminal command: `$ which qsub`

In case you have access to a multi-core computer you can follow the next steps to set up your own Grid Engine cluster and speed up your calculations:

```
# updated 02032017

# 1) go to http://arc.liv.ac.uk/downloads/SGE/releases ,
# create user 'sgeadmin' and download the latest binary packages
# (Debian-like here) matching your architecture (amd64 here):

wget -c http://arc.liv.ac.uk/downloads/SGE/releases/8.1.9/sge-common_8.1.9_all.deb
wget -c http://arc.liv.ac.uk/downloads/SGE/releases/8.1.9/sge_8.1.9_amd64.deb
wget -c http://arc.liv.ac.uk/downloads/SGE/releases/8.1.9/sge-dbg_8.1.9_amd64.deb

sudo useradd sgeadmin
sudo dpkg -i sge-common_8.1.9_all.deb
sudo dpkg -i sge_8.1.9_amd64.deb
sudo dpkg -i sge-dbg_8.1.9_amd64.deb
sudo apt-get install -f

# 2) set hostname to anything but localhost by editing /etc/hosts so that
# the first line is something like this (localhost or 127.0.x.x IPs not valid):
# 172.1.1.1    yourhost

# 3) install Grid Engine server with defaults except cluster name ('yourhost')
# and admin user name ('sgeadmin'):
sudo su
cd /opt/sge/
chown -R sgeadmin sge
chgrp -R sgeadmin sge
./install_qmaster

# 4) install Grid Engine client with all defaults:
./install_execd
exit

# 5) check the path to your sge binaries, which can be 'lx-amd64'
ls /opt/sge/bin

# 6) Set relevant environment variables in /etc/bash.bashrc [can also be named /etc/basrhc]
# or alternatively in ~/.bashrc for a given user
export SGE_ROOT=/opt/sge
export PATH=$PATH:"$SGE_ROOT/bin/lx-amd64"
```

```
# 7) Optionally configure default all.q queue:
qconf -mq all.q

# 8) Add your host to list of admitted hosts:
qconf -as yourhost
```

For cluster-based operations three bundled Perl scripts are invoked:

`_cluster_makeHomolog.pl`, `_cluster_makeInparalog.pl` and `_cluster_makeOrtholog.pl`.

It is also possible to invoke Pfam domain scanning from *get_homologues-est*. This option requires the bundled binary *hmmscan*, which is part of the [HMMER3](#) package, whose path is set in file *lib/phyTools.pm* (variable `$ENV{'EXE_HMMPFAM'}`). Should this binary not work in your system, a fresh install might be the solution, say in `/your/path/hmmer-3.1b2/`. In this case you'll have to edit file *lib/phyTools.pm* and modify the relevant:

```
if( ! defined($ENV{'EXE_HMMPFAM'}) )
{
$ENV{'EXE_HMMPFAM'} = '/your/path/hmmer-3.1b2/src/hmmscan --noali --acc --cut_ga ';
}
```

The Pfam HMM library is also required and the *install.pl* script should take care of it. However, you can manually download it from the appropriate [Pfam FTP site](#). This file needs to be decompressed, either in the default *db* folder or in any other location, and then it should be formatted with the program *hmmpress*, which is also part of the *HMMER3* package. A valid command sequence could be:

```
$ cd db;
$ wget ftp://ftp.sanger.ac.uk/pub/databases/Pfam/current_release/Pfam-A.hmm.gz .;
$ gunzip Pfam-A.hmm.gz;
$ /your/path/hmmer-3.1b2/src/hmmpress Pfam-A.hmm
```

Finally, you'll need to edit file *lib/phyTools.pm* and modify the relevant line to:

```
if( ! defined($ENV{"PFAMDB"}) ) { $ENV{"PFAMDB"} = "db/Pfam-A.hmm"; }
```

In order to reduce the memory footprint of *get_homologues-est* it is possible to take advantage of the [Berkeley_DB](#) database engine, which requires Perl core module [DB_File](#), which should be installed on all major Linux distributions.

The accompanying script *transcripts2cds.pl* should work out of the box, but the more efficient *transcripts2cdsCPP.pl* requires the installation of module [Inline::CPP](#), which in turn requires [Inline::C](#) and `g++`, the GNU C++ compiler. The installation of these modules is known to be troublesome in some systems, but the standard way should work in most cases:

```
$ yum -y install gcc-c++ perl-Inline-C perl-Inline-CPP # Redhat and derived distros

$ sudo apt-get -y install g++ # Ubuntu/Debian-based distros, and then cpan below

$ cpan -i Inline::C Inline::CPP # will require administrator privileges (sudo)
```

This script may optionally use [Diamond](#) instead of BLASTX. The bundled linux binary should work out of the box; in case the macOS binary does not work in your system you might have to re-compile it with:

```
cd bin/diamond-0.8.25/
build_simple.sh
cd ../../..
```

Similarly, in order to take full advantage of the accompanying script *parse_pangenome_matrix.pl*, particularly for option `-p`, the installation of module [GD](#) is recommended. An easy way to install them, provided that you have administrator privileges, is with help from the software manager of your Linux distribution (such as *synaptic/apt-get* in Ubuntu, *yum* in Fedora or *YaST* in openSUSE).

This can usually be done on the terminal as well, in different forms:

```
$ sudo apt-get -y install libgd-gd2-perl # Ubuntu/Debian-based distros
$ yum -y install perl-GD # Redhat and derived distros
$ zypper --assume-yes install perl-GD # SuSE
$ cpan -i GD # will require administrator privileges (sudo)
$ perl -MCPAN -e 'install GD' # will require administrator privileges (sudo)
```

The installation of perl-GD on macOSX systems is known to be [troublesome](#).

The accompanying scripts *compare_clusters.pl*, *plot_pancore_matrix.pl*, *parse_pangenome_matrix.pl*,

plot_matrix_heatmap.sh, *hcluster_matrix.sh* require the installation of the statistical software [R](#), which usually is listed by software managers in all major Linux distributions. In some cases (some [SuSE versions](#) and some [Redhat-like distros](#)) it will be necessary to add a repository to your package manager. R can be installed from the terminal:

```
$ sudo apt-get -y install r-base r-base-dev # Ubuntu/Debian-based distros
$ yum -y install R # RedHat and derived distros
$ zypper --assume-yes R-patched R-patched-devel # Suse
```

Please visit [CRAN](#) to download and install R on macOSX systems, which is straightforward.

In addition to R itself, *plot_matrix_heatmap.sh* and *hcluster_matrix.sh* require some R packages to run, which can be easily installed from the R command line with:

```
> install.packages(c("ape", "gplots", "cluster"), dependencies=TRUE)
```

Finally, the script *compare_clusters.pl* might require the installation of program PARS from the [PHYLIP suite](#), which should be already bundled with your copy of *get_homologues*.

3 User manual

This section describes the available options for the *get_homologues-est* software.

3.1 Input data

This program takes input sequences in FASTA format, which might be GZIP- or BZIP2-compressed, contained in a directory or folder containing several files with extension '.fna', which can have twin .faa files with translated amino acid sequences for the corresponding CDSs. File names matching the tag 'fcdna' are handled as full-length transcripts, and this information will be used downstream in order to estimate coverage. Global variable \$MINSEQLENGTH controls the minimum length of sequences to be considered; the default value is 20.

3.2 Program options

Typing `$./get_homologues-est.pl -h` on the terminal will show the basic options:

`-v` print version, credits and checks installation
`-d` directory with input FASTA files (`.fna` , optionally `.faa`), (use of pre-clustered sequences
1 per sample, or subdirectories (`subdir.clusters/subdir_`) ignores `-c`)
with pre-clustered sequences (`.faa/.fna`). Files matching
tag 'flcdna' are handled as full-length transcripts.
Allows for files to be added later.
Creates output folder named 'directory_est_homologues'

Optional parameters:

`-o` only run BLASTN/Pfam searches and exit (useful to pre-compute searches)
`-i` cluster redundant isoforms, including those that can be (min overlap, default: `-i 40`,
concatenated with no overhangs, and perform use `-i 0` to disable)
calculations with longest
`-c` report transcriptome composition analysis (follows order in `-I` file if enforced,
with `-t N` skips clusters occup<N [OMCL]
ignores `-r,-e`)
`-R` set random seed for genome composition analysis (optional, requires `-c`, example `-R 1234`)
`-s` save memory by using BerkeleyDB; default parsing stores
sequence hits in RAM
`-m` runmode [local|cluster] (default: `-m local`)
`-n` nb of threads for BLASTN/HMMER/MCL in 'local' runmode (default=2)
`-I` file with `.fna` files in `-d` to be included (takes all by default, requires `-d`)

Algorithms instead of default bidirectional best-hits (BDBH):

`-M` use orthoMCL algorithm (OMCL, PubMed=12952885)

Options that control sequence similarity searches:

`-C` min %coverage of shortest sequence in BLAST alignments (range [1-100], default: `-C 75`)
`-E` max E-value (default: `-E 1e-05` , max=0.01)
`-D` require equal Pfam domain composition (best with `-m cluster` or `-n threads`)
when defining similarity-based orthology
`-S` min %sequence identity in BLAST query/subj pairs (range [1-100], default: `-S 1` [BDBH|OMCL])
`-b` compile core-transcriptome with minimum BLAST searches (ignores `-c` [BDBH])

Options that control clustering:

`-t` report sequence clusters including at least `t` taxa (default: `t=numberOfTaxa`,
`t=0` reports all clusters [OMCL])
`-L` add redundant isoforms to clusters (optional, requires `-i`)
`-r` reference transcriptome `.fna` file (by default takes file with
least sequences; with BDBH sets
first taxa to start adding genes)
`-e` exclude clusters with inparalogues (by default inparalogues are
included)
`-F` orthoMCL inflation value (range [1-5], default: `-F 1.5` [OMCL])
`-A` calculate average identity of clustered sequences, (optional, creates tab-separated matrix,
uses blastn results [OMCL])
`-z` add soft-core to genome composition analysis (optional, requires `-c` [OMCL])

The only required option is `-d`, which indicates an input folder, as seen in section 3.1. It is important to remark that in principle only files with extensions `.fna` / `.fa` / `.fasta` and optionally `.faa` are considered when parsing the `-d` directory. By using `.faa` input files protein sequences can be used to scan Pfam domains and included in output clusters.

The use of an input folder or directory (`-d`) is recommended as it allows for new files to be added there in the future, reducing the computing required for updated analyses. For instance, if a user does a first analysis with 5 input genomes

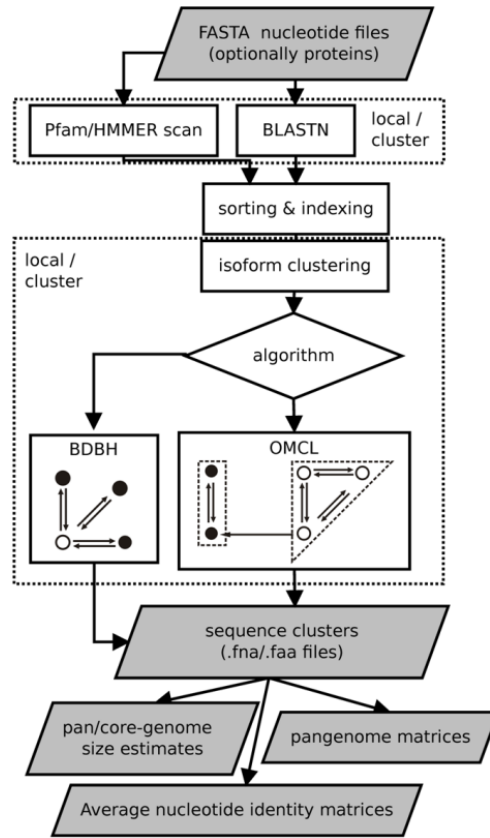


Figure 1: Flowchart of get.homologues-est.

name	option	
BDBH	default	Starting from a reference genome, keep adding genomes stepwise while storing the sequence clusters that result of merging the latest bidirectional best hits.
OMCL	-M	OrthoMCL v1.4, uses the Markov Cluster Algorithm to group sequences, with inflation (-F) controlling cluster granularity, as described in PubMed= 12952885 .

Table 2: List of available clustering algorithms. Note that the COG triangles algorithm is not supported.

today, it is possible to check how the resulting clusters would change when adding an extra 10 genomes tomorrow, by copying these new 10 .fna input files to the pre-existing -d folder, so that all previous BLASTN searches are re-used.

All remaining flags are options that can modify the default behavior of the program, which is to use the bidirectional best hit algorithm (BDBH) in order to compile clusters of potential orthologous DNA sequences, taking the smallest genome as a reference. By default nucleotide sequences are used to guide the clustering, thus relying on BLASTN searches.

Perhaps the most important optional parameter would be the choice of clustering algorithm (Table 2):

The remaining options are now reviewed:

- Apart from showing the credits, option -v can be helpful after installation, for it prints the enabled features of the program.
- -o is ideally used to submit to a computer cluster the required BLAST (and Pfam) searches, preparing a job for posterior analysis on a single computer.
- -i can be used to filter out short, redundant isoforms which overlap, with no overhangs, for a minimum length. By default this is set to \$MINREDOVERLAP=40 as in PubMed=[12651724](#). This EST-specific feature can be turned off by setting -i 0. Redundant isoforms will not be output unless -L is set.



Figure 2: Redundant isoforms (dashed) are optionally removed from an input sequence set if they overlap a longer sequence over a length ≥ 40 (A) or when they are completely matched (B). In either case a 100% sequence identity is required. By calling option -L all redundant isoforms are included in the output.

- -c is used to request a pan- and core-genome analysis of the input sequences, which will be output as tab-separated files. The number of samples for the genome composition analysis is set to 20 by default, but this can be edited at the header of `get_homologues-est.pl` (check the `$NOFSAMPLESREPORT` variable). As *get_homologues-est* is meant to be used primarily for the study of transcripts/CDSs of the same species, it uses appropriate thresholds to define new accessory genes (`$MIN_PERSEQID_HOM_EST=70`, `$MIN_COVERAGE_HOM_EST=50`), which mean that genes/transcripts added to the pool must be $\geq 70\%$ identical in sequence to any previous sequence with cover $\geq 50\%$ in order to be called homologous; otherwise they are handled as novel sequences. This low coverage is set in order to allow transcripts with retained/unprocessed introns to be matched. The default identity value was chosen to match the fact that BLASTN::megablast hardly reports hits with lower identities in our tests with barley transcripts and *A.thaliana* CDS sequences. Note that these default values are different to those set in *get_homologues* for peptide sequences. When combined with flag -t (see below), the composition analysis will disregard clusters reported in a selected number of strains. This feature can be used to filter out singletons or artifacts which might arise from *de novo* assembled transcriptomes.




Figure 3: Histograms of % identity reported by BLASTN among *Arabidopsis thaliana* CDS sequences (left) and *Hordeum vulgare* (right) transcripts. Note that the default BLASTN algorithm (megablast) hardly reports alignments with identities $< 70\%$. Plots were computed from 51,110,547 and 70,653,179 alignments, respectively.

- -R takes a number that will be used to seed the random generator used with option -c. By using the same seed in different -c runs the user ensures that genomes are sampled in the same order.

- `-s` can be used to reduce the memory footprint, provided that the Perl module [BerkeleyDB](#) is in place. This option usually makes *get_homologues-est* slower, but for very large datasets or in machines with little memory resources this might be the only way to complete a job.
- `-m` allows the choice of runmode, which can be either `-m local` (the default) or `-m cluster`. In the second case global variable `$SGEPATH` might need to be appropriately set, as explained in [manual.get_homologues.pdf](#), as well as `$QUEUESETTINGS`, that specifies for instance a particular queue name for your cluster jobs.
- `-n` sets the number of threads/CPU's to dedicate to each BLAST/HMMER/mcl job run locally, which by default is 2.
- `-I list_file.txt` allows the user to restrict a *get_homologues-est* job to a subset of FASTA files included in the input `-d` folder. This flag can be used in conjunction with `-c` to control the order in which genomes are considered during pan- and core-transcriptome analyses. Taking the `sample_RNAseq` folder, a valid `list_file.txt` could contain these lines:

```
Esterel.trinity.fna.bz2
Franka.trinity.fna.bz2
```

- option `-C` sets the minimum percentage of coverage required to call two sequences best hits. As EST/transcripts are frequently truncated, by default coverage is calculated with respect to the shortest sequence in the pair, unless both of them come from a full-length collection (see [3.1](#)).




Figure 4: Coverage illustrated with the alignment of sequence 'query' to two aligned fragments of sequence 'subject'. L_q and L_s are the lengths of both sequences, and $s1, e1, s2, e2$ and L_q are alignment coordinates.

- `-E` sets the maximum expectation value (E-value) for BLASTN alignments. This value is by default set to $1e-05$.
- `-D` is an extra restriction for calling best hits, that should have identical Pfam domain compositions. Note that this option requires scanning all input sequences for Pfam domains, and this task requires extra computing time, ideally on a computer cluster (`-m cluster`). While for BDBH domain filtering is done at the time bidirectional best hits are called, this processing step is performed only after the standard OMCL algorithms have completed, to preserve the algorithm features.
- `-S` can be passed to require a minimum % sequence identity for two sequences to be called best hits. The default value is set to 95%, as in PubMed=[21572440](#).
- `-b` reduces the number of pairwise BLAST searches performed while compiling core-genomes with algorithm BDBH, reducing considerably memory and run-time requirements (for G genomes, $3G$ searches are launched instead of the default G^2). It comes at the cost of being less exhaustive in finding inparalogues, but in our bacterial benchmarks this potential, undesired effect was negligible.
- `-t` is used to control which sequence clusters should be reported. By default only clusters which include at least one sequence per genome are output. However, a value of `-t 2` would report all clusters containing sequences from at least 2 taxa. A special case is `-t 0`, which will report all clusters found, even those with sequences from a single genome.
- `-r` allows the choice of any input sequence set (of course included in `-d` folder) as the reference, instead of the default smaller one. If possible, resulting clusters are named using CDS/transcript names from this genome, which can be used to select well annotated species for this purpose.

- `-e` excludes clusters with inparalogues, defined as sequences with best hits in its own genome. This option might be helpful to rule out clusters including several sequences from the same species, which might be of interest for users employing these clusters for primer design, for instance.
- `-F` is the inflation value that governs Markov Clustering in OMCL runs, as explained in PubMed=[12952885](#). As a rule of thumb, low inflation values (`-F 1`) result in the inclusion of more sequences in fewer groups, whilst large values produce more, smaller clusters (`-F 4`).
- `-A` tells the program to produce a tab-separated file with average % sequence identity values among pairs of genomes, computed from sequences in the final set of clusters (see also option `-t`). By default these identities are derived from BLASTN alignments, and hence correspond to nucleotide sequence identities, to produce genomic average nucleotide sequence identities (ANI).
- `-z` can be called when performing a genome composition analysis with clustering algorithm OMCL. In addition to the core- and pan-genome tab-separated files mentioned earlier (see option `-c`), this flag requests a soft-core report, considering all sequence clusters present in a fraction of genomes defined by global variable `$SOFTCOREFRACTION`, with a default value of 0.95. This choice produces a composition report more robust to assembly or annotation errors than the core-genome.

3.3 Accompanying scripts

The following Perl and shell scripts are included in each release to assist in the interpretation of results generated by *get_homologues-est.pl*. See examples of use in *manual_get_homologues.pdf* [manual_get_homologues.pdf](#):

- *compare_clusters.pl* primarily calculates the intersection between cluster sets, which can be used to select clusters supported by different algorithms or settings. This script can also produce pangenome matrices and Venn diagrams.
- *parse_pangenome_matrix.pl* is a script that can be used to analyze pan-genome sets, in order to find transcripts/genes present in a group A of strains which are absent in set B. This script can also be used for calculating and plotting cloud, shell and core genome compartments.
- *make_nr_pangenome_matrix.pl* is provided to post-process pangenome matrices in case the user wishes to remove redundant clusters, using either nucleotide or protein sequence identity cut-offs.
- *plot_pancore_matrix.pl*, a Perl script to plot pan/soft/core-genome sampling results and to fit regression curves with help from [R](#) functions.
- *check_BDBHs.pl* is a script that can be used, after a previous *get_homologues-est* run, to find out the bidirectional best hits of a sequence identifier chosen by the user. It can also retrieve the Pfam annotations of a sequence and its reciprocal best hits.
- *add_pancore_matrices.pl* can be used to add pan/core-matrices produced by previous *get_homologues-est -c -R* runs on the same set of genomes, with the aim of combining clusters.
- *annotate_cluster.pl* can be used to retrieve a multiple alignment view of the supporting local BLAST alignments of the sequences in the cluster, and to annotate any encoded Pfam domain.
- *plot_matrix_heatmap.sh* calculates ordered heatmaps with attached row and column dendrograms from tab-separated numeric matrices, which can be presence/absence pangenomic matrices or similarity / identity matrices as those produced by *get_homologues-est* with flag *-A*.
- *hcluster_matrix.sh* generates a distance matrix out of a tab-separated numeric matrix, which is then used to call [R](#) functions `hclust()` and `heatmap.2()` in order to produce a heatmap.
- *pfam_enrich.pl* calculates the enrichment of a set of sequence clusters in terms of Pfam domains, by using [Fisher's exact test](#).

Apart from these, auxiliary *transcripts2cds.pl* script is bundled to assist in the analysis of transcripts. In particular, this script can be used to annotate potential Open Reading Frames (ORFs) contained within raw transcripts, which might be truncated or contain introns. This script uses [TransDecoder](#), BLASTX and SWISSPROT, which should be installed by running: *./install.pl*

This program supports the following options:

usage: *./transcripts2cds.pl* [options] <input FASTA file(s) with transcript nucleotide sequences>

```
-h this message
-p check only 'plus' strand                (optional, default both strands)
-l min length for CDS                      (optional, default=50 amino acid residues)
-g genetic code to use during translation  (optional, default=1, example: -g 11)
-d run blastx against selected protein FASTA database file (default=swissprot, example: -d db.faa)
-E max E-value during blastx search        (default=1e-05)
-n number of threads for BLASTX jobs       (default=2)
-X use diamond instead of blastx           (optional, much faster for many sequences)

-G show available genetic codes and exit
```

The main output of this script are two files, as shown in section [4.1](#), which contain inferred nucleotide and peptide CDS sequences. These FASTA files contain in each header the evidence supporting each called CDS, which can be blastx, transdecoder or a combination of both, giving precedence to blastx in case of conflict. Note that we have observed that the output of TransDecoder might change if a single sequence is analyzed alone, in contrast to the

analysis of a batch of sequences. The next table shows the rules and evidence codes used by this script in order to call CDS sequences by merging BLASTX (1) and TransDecoder (2) predictions. The rules are mutually exclusive and are tested hierarchically from top to bottom. Sequences from 1 and 2 with less than 90 consecutive matches (30 amino acid residues) are considered to be non-overlapping (last rule). Note that the occurrence of mismatches are checked as a control:

graphical summary	evidence	description
1-----	blastx	no transdecoder
2-----	transdecoder	no blastx
1----- 2-----	blastx.transdecoder	inferred CDS overlap with no mismatches and are concatenated
1----- 2-----	transdecoder.blastx	inferred CDS overlap with no mismatches and are concatenated
1----- 2-----	blastx<transdecoder	blastx CDS includes transdecoder CDS
1----- 2-----	transdecoder<blastx	transdecoder CDS includes blastx CDS
1-----C-- 2---T-----	blastx-mismatches	blastx CDS is returned as sequences have mismatches
1----- 2---	blastx-noover	blastx CDS is returned as transdecoder CDS does not overlap

Our benchmarks suggest that 78 to 92% of deduced CDS sequences match the correct peptide sequences. :

evidence	<i>Arabidopsis thaliana</i> [Col-0]	n	<i>Hordeum vulgare</i> [Haruna Nijo]	n
blastx	0.787	960	0.654	1,657
transdecoder	0.914	8,194	0.662	9,026
blastx.transdecoder	0.930	4,678	0.843	5,939
transdecoder.blastx	0.959	15,700	0.859	8,903
blastx<transdecoder	0.620	324	0.674	218
transdecoder<blastx	0.966	6,581	0.872	11,999
blastx-mismatches	0	1		0
blastx-noover	0.232	835	0.426	2,211
overall	0.923		0.783	

Table 3: Fraction of correct peptide sequences in deduced CDS obtained by combining BLASTX and TransDecoder.

The results obtained with DIAMOND instead of BLASTX are very similar:

evidence	<i>Arabidopsis thaliana</i> [Col-0]	n	<i>Hordeum vulgare</i> [Haruna Nijo]	n
blastx	0.800	929	0.655	1,598
transdecoder	0.914	8,166	0.663	8,980
blastx.transdecoder	0.929	4,685	0.844	5,951
transdecoder.blastx	0.958	15,698	0.859	8,890
blastx<transdecoder	0.615	325	0.671	216
transdecoder<blastx	0.967	6,583	0.872	11,999
blastx-mismatches		0		0
blastx-noover	0.270	833	0.452	2,190

Table 4: Fraction of correct peptide sequences in deduced CDS obtained by combining DIAMOND and TransDecoder.

4 A few examples

This section presents typical ways of running *get_homologues-est.pl* and the accompanying scripts with provided sample input data. Please check file [manual_get_homologues.pdf](#) for more examples, particularly for the auxiliary scripts, which are not explained in this document.

4.1 Extracting coding sequences (CDS) from transcripts

This example takes the provided sample file `sample_transcripts.fna` to demonstrate how to annotate coding sequences contained in those sequences by calling `transcripts2cds.pl`. Note that `transcripts2cdsCPP.pl` is significantly faster, but requires an optional Perl module (see 2.3).

This is an optional pre-processing step which you might not want to do, as the software should be able to properly handle any nucleotides sequences suitable for BLASTN. However, coding sequences have the advantage that can be translated to amino acids and thus used to scan Pfam domains further down in the analysis (see option -D).

A simple command would be, which will discard sequences less than 50b long, and will aligned them to SWISSPROT proteins in order to annotate coding regions. In case of overlap, Transdecoder-defined and BLASTX-based coding regions are combined provided that a \$MINCONOVERLAP=90 overlap, with no mismatches, is found; otherwise the latter are given higher priority:

```
./transcripts2cdsCPP.pl -n 10 sample_transcripts.fna
```

The output should look like this (contained in file `sample_transcripts_output.txt`):

```
# ./transcripts2cdsCPP.pl -p 0 -m -d /path/get_homs-est/db/swissprot -E 1e-05 -l 50 -g 1 -n 10 -X 0
# input files(s):
# sample_transcripts.fna

## processing file sample_transcripts.fna ...
# running transdecoder...

# parsing transdecoder output (_sample_transcripts.fna_l50.transdecoder.cds.gz) ...
# running blastx...
# parsing blastx output (_sample_transcripts.fna_E1e-05.blastx.gz) ...
# calculating consensus sequences ...
# input transcripts = 9
# transcripts with ORFs = 7
# transcripts with no ORFs = 2
# output files: sample_transcripts.fna_l50_E1e-05.transcript.fna ,
# sample_transcripts.fna_l50_E1e-05.cds.fna ,
# sample_transcripts.fna_l50_E1e-05.cds.faa ,
# sample_transcripts.fna_l50_E1e-05.noORF.fna
```

The resulting CDS files can then be analyzed with *get_homologues-est.pl*.

Apart from the listed output files, which include translated protein sequences, temporary files are stored in the working directory, which of course can be removed, but will be re-used if the same job is re-run later, such as `_sample_transcripts.fna_E1e-05.blastx.gz`,

`_sample_transcripts.fna_150.transdecoder.cds.gz` or
`_sample_transcripts.fna_150.transdecoder.pep.gz`.

By default the script uses BLASTX (in combination with Transdecoder), which might take quite some time to process large numbers of sequences. For this reason the DIAMOND algorithm is also available (upon calling option -X), which in our [benchmarks](#) showed comparable performance and was several orders of magnitude faster when using multiple CPU cores.

CDS sequences can be deduced for a collection of transcriptomes and put in the same folder, so that they can all be analyzed together with *get_homologues-est.pl*. Such files support calling option -D, which will annotate Pfam domains contained in those sequences, and can then also be used to calculate enrichment as explained in [manual_get_homologues.pdf](#).

4.2 Clustering orthologous transcripts from FASTA files, one per strain

This example takes the sample input folder `sample_transcripts_fasta`, which contains automatically assembled transcripts (Trinity) of three *Hordeum vulgare* strains (barley), plus a set of full-length cDNA collection of cultivar *Haruna Nijo*, to show to produce clusters of transcripts.

The next command uses the OMCL algorithm to cluster sequences, produces a composition report, including the soft-core, and finally computes an Average Nucleotide Identity matrix on the produced clusters. Note that redundant isoforms are filtered, keeping only the longest one (you can turn this feature off with `-i 0`):

```
$ ./get_homologues-est.pl -d sample_transcripts_fasta -M -c -z -A
```

The output should look like this (contained in file `sample_transcripts_output.txt`):

```
# results_directory=/path/sample_transcripts_fasta_est_homologues
# parameters: MAXEVALUEBLASTSEARCH=0.01 MAXPFAMSEQS=250 BATCHSIZE=100

# checking input files...
# Esterel.trinity.fna.bz2 5892 median length = 506
# Franka.trinity.fna.bz2 6036 median length = 523
# Hs_Turkey-19-24.trinity.fna.bz2 6204 median length = 476
# flcdnas_Hnijo.fna.gz 28620 [full length sequences] median length = 1504

# 4 genomes, 46752 sequences

# taxa considered = 4 sequences = 46752 residues = 63954041

# mask=Esterel_alltaxa_algOMCL_e0_ (_algOMCL)
[...]

# re-using previous isoform clusters
# 42 sequences
# 65 sequences
# 61 sequences
# 2379 sequences

# creating indexes, this might take some time (lines=2.08e+05) ...

# construct_taxa_indexes: number of taxa found = 4
# number of file addresses/BLAST queries = 4.4e+04

# genome composition report (samples=20,permutations=24,seed=0)
# genomic composition parameters: MIN_PERSEQID_HOM=70 MIN_COVERAGE_HOM=50 SOFTCOREFRACTION=0.95
[...]

# file=sample_transcripts_fasta_est_homologues/core_genome_algOMCL.tab
genomes mean stddev | samples
0 8559 6614 | 4665 4665 4665 ...
1 1113 737 | 496 432 2007 ...
2 255 101 | 84 308 347 ...
3 66 0 | 66 66 66 ...

# file=sample_transcripts_fasta_est_homologues/soft-core_genome_algOMCL.tab
genomes mean stddev | samples
0 8559 6614 | 4665 4665 4665 ...
1 3491 2311 | 2428 2195 8108 ...
2 2170 1017 | 765 3460 2145 ...
3 645 101 | 816 592 553 ...
```

```
# looking for valid sequence clusters (n_of_taxa=4)...

# number_of_clusters = 66
# cluster_list = sample_transcripts_fasta_est_homologues/Esterel_alltaxa_algOMCL_e0_.cluster_list
# cluster_directory = sample_transcripts_fasta_est_homologues/Esterel_alltaxa_algOMCL_e0_

# average_nucleotide_identity_matrix_file = # [...] /Esterel_alltaxa_algOMCL_e0_Avg_identity.tab
```

Notice that both core and soft-core sampling experiments are reported, considering sequences found in all strains and in 95% strains, respectively. The produced Average Nucleotide Identity matrix looks like this:

genomes	Esterel	Franka	HsTurkey	flcdnasHnijo
Esterel	100	98.29	98.04	99.33
Franka	98.29	100	98.25	98.90
HsTurkey	98.04	98.25	100	98.41
flcdnasHnijo	98.33	98.90	98.41	100

Provided that optional R modules described in [manual_get_homologues.pdf](#) are installed, this matrix can be plotted with the following command:

```
./plot_matrix_heatmap.sh -i sample_[...] /Esterel_alltaxa_algOMCL_e0_Avg_identity.tab \
-t "clusters=66" -k "Average Nucleotide Identity" -o pdf -m 28 -v 35 -H 9 -W 10
```

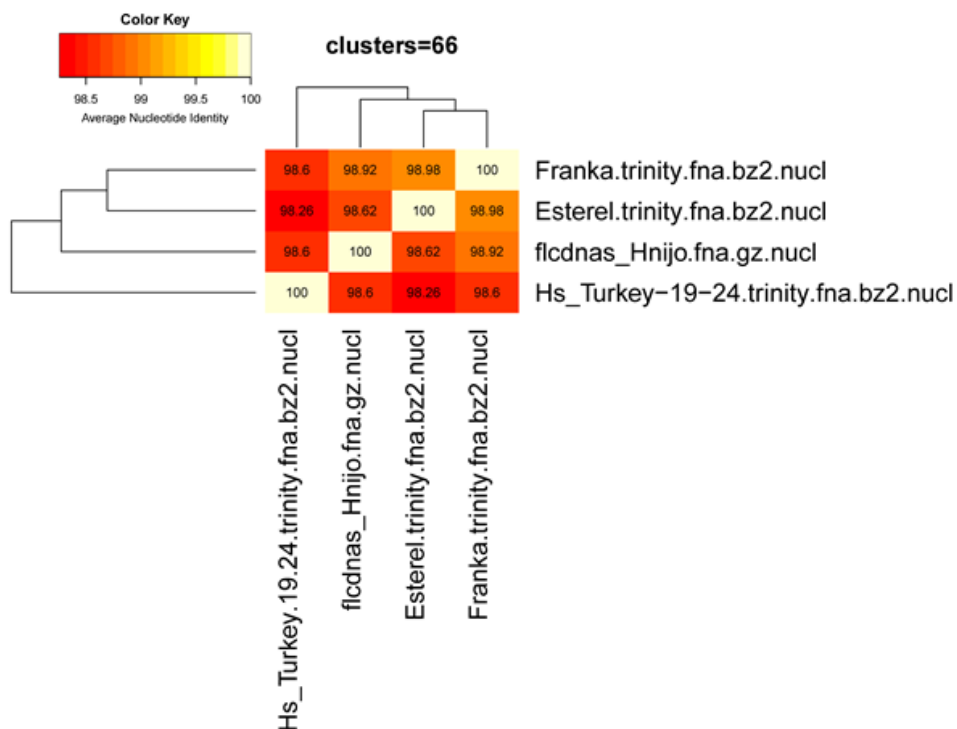


Figure 5: Heatmap of Average Nucleotide Identity.

If the previous command is changed by adding option `-t -2` only transcripts present in at least two strains will be considered, which are output in folder:

`sample_transcripts_fasta_est_homologues/Esterel_2taxa_algOMCL_e0_`

This second command produces a significantly different pan-genome composition matrix, which changes from:

```
# file=sample_transcripts_fasta_est_homologues/pan_genome_algOMCL.tab
genomes mean stddev | samples
0 8559 6614 | 4665 4665 4665 ...
1 14830 6425 | 8937 9002 21292 ...
2 21384 4866 | 13004 24283 23358 ...
3 26380 468 | 27019 26209 25652 ...

to

# file=sample_transcripts_fasta_est_homologues/pan_genome_2taxa_algOMCL.tab
genomes mean stddev | samples
0 2860 1172 | 2262 2262 2262 ...
1 4270 490 | 4110 3828 4196 ...
2 4953 424 | 5475 4767 4294 ...
3 4954 424 | 5475 4768 4296 ...
```

Both matrices can be plotted with script `plot_pancore_matrix.pl`, with a command such as:

```
./plot_pancore_matrix.pl -i sample_transcripts_fasta_est_homologues/pan_genome_algOMCL.tab -f pan
```




Figure 6: Pan-transcriptome size estimates (`-t 0`, left) and (`-t 2`, right) based on random samples of 4 transcriptome sets. As the left example illustrates, four strains are usually not enough to fit a Tettelin-like function.

The next figure shows a similar analysis but now using genomic data instead of transcript sets. The example shows pan-genome size estimates of Whole Genome Sequence assemblies of 19 *Arabidopsis thaliana* ecotypes, downloaded from <http://mus.well.ox.ac.uk/19genomes/sequences/CDS> and described in PubMed=21874022.

Script `plot_pancore_matrix.pl` can also be called with flag `-a`:

```
./plot_pancore_matrix.pl -i sample_transcripts_fasta_est_homologues/pan_genome_algOMCL.tab \
-f pan -a snapshots+
```

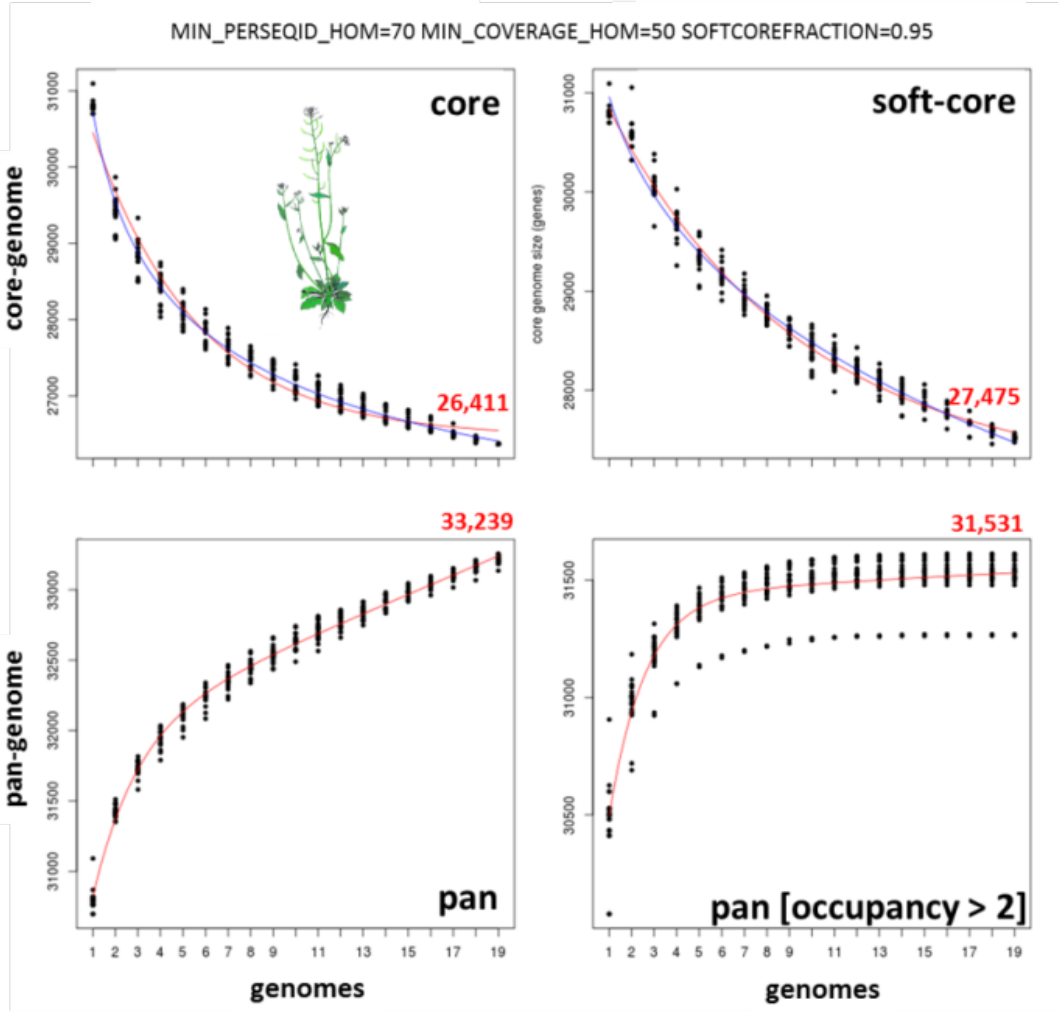


Figure 7: Core-genome, soft-core-genome and pan-genome CDS composition analysis of WGS assemblies of 19 *A.thaliana* ecotypes. Note that the pan-genome simulation was done with all clusters (left) and with all clusters found in at least three genomes (right), illustrating the effect of option -t 3, which might be useful to remove low confidence sequences. Red numbers correspond to fitted values generated by *plot-pancore_matrix.pl*.

This will create and store a in folder `snapshots/` a series of GIF images that can be used to animate pan-genome simulations. The next Figure show some of this snapshots:




Figure 8: Four snapshots of the pan-genome simulation carried out in the previous figure, generated by `plot_pancore_matrix.pl`.

4.3 Producing a nucleotide-based pangenome matrix

The clusters obtained in the previous section with option `-t 2` can be used to compile a pangenome matrix without singletons with this command:

```
./compare_clusters.pl -d sample_[...]/Esterel_2taxa_algOMCL_e0_ -o outdir -n -m
```

```
# number of input cluster directories = 1
```

```
# parsing clusters in sample_transcripts_fasta_est_homologues/Esterel_2taxa_algOMCL_e0_ ...  
# cluster_list in place, will parse it (sample_[...]/Esterel_2taxa_algOMCL_e0_.cluster_list)  
# number of clusters = 5243
```

```
# intersection output directory: outdir
```

```
# intersection size = 5243 clusters
```

```
# intersection list = outdir/intersection_t0.cluster_list
```

```
# pangenome_file = outdir/pangenome_matrix_t0.tab
```

```
# pangenome_phylip file = outdir/pangenome_matrix_t0.phylip
```

If the optional R modules described in [manual_get_homologues.pdf](#) are installed, such a pangenome matrix can be used to hierarchically cluster strains with this command:

```
./hcluster_matrix.sh -i outdir/pangenome_matrix_t0.tab
```




Figure 9: Hierarchical grouping of strains based on pangenome matrix.

4.4 Estimating protein domain enrichment of some sequence clusters

This example uses data from the barley benchmark, the `test_barley/` folder, which contains instructions to download sequences from: <http://floresta.eead.csic.es/plant-pan-genomes>

After completing the downloads, the folder will contain FASTA files with nucleotide sequences of 14 de-novo assembled transcriptomes and transcripts/cDNA sequences annotated in reference accessions Morex and Haruna Nijo. CDS can be extracted as explained in Section 5 and then Pfam domains can be annotated as follows:

```
$ ./get_homologues-est.pl -d cds -D -o -m cluster
```

These annotations will serve to calculate background domain frequencies.

Once this is completed, we can compute "control" clusters with this command:

```
$ ./get_homologues-est.pl -d cds -M -m cluster
```

which we will then place in a folder called `clusters_cds`:

```
compare_clusters.pl -d cds_est_homologues/Alexis_0taxa_algOMCL_e0_ \
-o clusters_cds -m -n+
```

In order to call accessory sequences with more confidence we will use only non-cloud clusters (*occupancy* > 2), which we do with this command:

```
$ ./get_homologues-est.pl -d cds -M -t 3 -m cluster
```

The output should include the next lines:

```
# number_of_clusters = 34248
# cluster_list = cds_est_homologues/Alexis_3taxa_algOMCL_e0_.cluster_list
# cluster_directory = cds_est_homologues/Alexis_3taxa_algOMCL_e0_
```

We should be now in position to compile the pan-genome matrix corresponding to these clusters:

```
./compare_clusters.pl -d cds_est_homologues/Alexis_3taxa_algOMCL_e0_ \
-o clusters_cds_t3 -m -n
```

which should produce:

```
# number of clusters = 34248

# intersection output directory: clusters_cds_t3
# intersection size = 34248 clusters

# intersection list = clusters_cds_t3/intersection_t0.cluster_list

# pangenome_file = clusters_cds_t3/pangenome_matrix_t0.tab
# pangenome_phylip file = clusters_cds_t3/pangenome_matrix_t0.phylip
```

We should now interrogate the pan-genome matrix, for instance looking for clusters found in one genotype (A) but not in others (B):

```
./parse_pangenome_matrix.pl -m clusters_cds_t3/pangenome_matrix_t0.tab \
-A cds/SBCC073.list -B cds/ref.list -g
```

You should obtain a list of 4348 accessory clusters:

```
# matrix contains 34248 clusters and 16 taxa

# taxa included in group A = 1

# taxa included in group B = 2

# finding genes present in A which are absent in B ...
# file with genes present in set A and absent in B (4348):
clusters_cds_t3/pangenome_matrix_t0__pangenomes_list.txt
```


Finally, we will now estimate whether these clusters are enriched in any Pfam domain, producing also a single FASTA file with the tested sequences:

```
./pfam_enrich.pl -d cds_est_homologues -c clusters_cds -n -t greater \
-x clusters_cds_t3/pangenome_matrix_t0__pangenes_list.txt -e -p 0.05 \
-r SBCC073 -f SBCC073_accessory.fna
```

The output should be:

```
# 39400 sequences extracted from 113222 clusters

# total experiment sequence ids = 4818
# total control      sequence ids = 39400

# parse_Pfam_freqs: set1 = 562 Pfams set2 = 3718 Pfams

# created FASTA file: SBCC073_accessory.fna

# sequences=4818 mean length=353.8 , seqs/cluster=1.11

# fisher exact test type: 'greater'
# multi-testing p-value adjustment: fdr
# adjusted p-value threshold: 1

# total annotated domains: experiment=1243 control=19192

#PfamID counts(exp) counts(ctr) freq(exp) freq(ctr) p-value p-value(adj) description
PF00009 0 20 0.000e+00 1.042e-03 1.000e+00 1.000e+00 Elongation factor Tu GTP binding domain
PF00010 0 32 0.000e+00 1.667e-03 1.000e+00 1.000e+00 Helix-loop-helix DNA-binding domain
...
PF00665 13 31 1.046e-02 1.615e-03 1.418e-06 1.318e-03 Integrase core domain
PF07727 28 61 2.253e-02 3.178e-03 3.033e-13 1.128e-09 Reverse transcriptase (RNA-dep DNA pol)
PF00931 44 201 3.540e-02 1.047e-02 1.750e-10 3.253e-07 NB-ARC domain
PF13976 14 19 1.126e-02 9.900e-04 2.744e-09 3.401e-06 GAG-pre-integrase domain
```

4.5 Making and annotating a non-redundant pangenome matrix

The script `make_nr_pangenome_matrix.pl` produces a non-redundant pangenome matrix by comparing all clusters to each other, taking the median sequence in each cluster. By default nucleotide sequences are compared, but if the original input of *get_homologues-est* comprised both DNA and protein sequences, the user can also choose peptide sequences to compute redundancy, which probably make more sense in terms of protein function. On the contrary, it would seem more appropriate to use DNA sequences to measure diversity.

In this example a DNA-based non-redundant pangenome matrix is computed with BLASTN assuming that sequences might be truncated (option `-e`) and using 10 processor cores and a coverage cutoff of 50%:

```
./make_nr_pangenome_matrix.pl -m outdir/pangenome_matrix_t0.tab -n 10 -e -C 50

# input matrix contains 5243 clusters and 4 taxa

# filtering clusters ...
# 5243 clusters with taxa >= 1 and sequence length >= 0

# sorting clusters and extracting median sequence ...

# running makeblastdb with outdir/pangenome_matrix_t0_nr_t1_l0_e1_C50_S90.fna

# parsing blast result! (outdir/pangenome_matrix_t0_nr_t1_l0_e1_C50_S90.blast , 0.34MB)
# parsing file finished

# 5210 non-redundant clusters
# created: outdir/pangenome_matrix_t0_nr_t1_l0_e1_C50_S90.fna

# printing nr pangenome matrix ...
# created: outdir/pangenome_matrix_t0_nr_t1_l0_e1_C50_S90.tab
```

Note that the previous command can be modified to match external reference sequences, for instance from [Swissprot](#), or pre-computed clusters, such as groups of orthologous sequences, so that the resulting matrix contains cross-references to those external clusters, and their annotations. In either case, both input clusters and reference sequences must be of the same type: either nucleotides or peptides.

The next example shows how a set of clusters produced by *get_homologues-est* can be matched to some nucleotide reference sequences, in this case annotated rice cDNAs:

```
./make_nr_pangenome_matrix.pl -m outdir/pangenome_matrix_t0.tab -n 10 -e -C 50 -f oryza.fna
This is the produced output:
```

```
# input matrix contains 5243 clusters and 4 taxa

# filtering clusters ...
# 5243 clusters with taxa >= 1 and sequence length >= 0

# sorting clusters and extracting median sequence ...
# re-using previous BLAST output outdir/pangenome_matrix_t0_nr_t1_l0_e1_C50_S90.blast

# parsing blast result! (outdir/pangenome_matrix_t0_nr_t1_l0_e1_C50_S90.blast , 0.34MB)
# parsing file finished

# 5210 non-redundant clusters
# created: outdir/pangenome_matrix_t0_nr_t1_l0_e1_C50_S90.fna

# 66339 reference sequences parsed in oryza.fna

# parsing blast result! (outdir/pangenome_matrix_t0_nr_t1_l0_e1_C50_S90_ref.blast , 0.37MB)
# parsing file finished
```

```
# matching nr clusters to reference (%alignment coverage cutoff=50) ...

# printing nr pangenome matrix ...
# created: outdir/pangenome_matrix_t0_nr_t1_l0_e1_C50_S90_ref_c50_s50.tab

# NOTE: matrix can be transposed for your convenience with:
```

```
perl -F'\t' -ane '$r++;for(1 .. @F){$m[$r][$_]=$F[$_-1]}; \
  $mx=@F;END{for(1 .. $mx){for $t(1 .. $r){print"$m[$t][$_]\\t"}print"\\n"}}' \
  outdir/pangenome_matrix_t0_nr_t1_l0_e1_C50_S90_ref_c50_s50.tab
```

The suggested perl command can be invoked to tranpose the matrix, which now contains rows such as these:

```
non-redundant Franka.bz2.nucl Esterel.bz2.nucl flcdnas_Hnijo.gz.nucl ... redundant reference
1_TR2804-c0_g1_i1.fna 1 1 0 0 NA LOC_0s09g07300.1 cDNA|BIG, putative, expressed
2_TR1554-c0_g1_i1.fna 0 2 1 0 NA LOC_0s03g53280.1 cDNA|WD domain containing protein
6_TR3918-c0_g1_i1.fna 0 1 1 0 NA NA
...
```

Pangenome matrices with more than 4 taxa can be plotted with help from script *parse_pangenome_matrix.pl*, as explained in [manual_get_homologues.pdf](#).

4.6 Annotating a sequence cluster

After analyzing pan-genome or pan-transcriptome clusters it might be interesting to find out what kind of proteins they encode, or we might just want to double-check the BLAST matches that support a produced cluster. The script *annotate_cluster.pl* does just that, and be used with both nucleotide and peptide clusters. It can be called like this:

```
./annotate_cluster.pl -f outdir/1004_TR425-c0_g2_i1.fna -o 1004_TR425-c0_g2_i1.aln.fna -D
```

And will produce this output:

```
# DEFBLASTNTASK=megablast DEFEVALUE=10
# MINBLUNTBLOCK=100 MAXSEQNAMELEN=60
# MAXMISMCOLLAP=0 MAXGAPSCOLLAP=2

# ./annotate_cluster.pl -f outdir/1004_TR425-c0_g2_i1.fna -r
# -o 1004_TR425-c0_g2_i1.aln.fna -P 1 -b 0 -D 1 -c 0

# total sequences: 3 taxa: 2

# Pfam domains: PF10602,PF01399,
# Pfam annotation: 26S proteasome subunit RPN7;PCI domain;
# aligned sequences: 3 width: 1595
# taxa included in alignment: 2
## Esterel.trinity 2
## flcdnas_Hnijo 1

# alignment file: 1004_TR425-c0_g2_i1.aln.fna
```

If option `-b` is enforced a blunt-end alignment is produced, which might be useful for further analyses. In either case, the produced FASTA alignment file will contain sequence variants and Pfam domains in each header, in addition to the relevant BLAST scores:

>TR425|c0_g2_i1_[Estere1.trinity.fna.bz2] bits E-value N qy ht 1:1595 SNPs:123,1169,.. Pfam:..
CCTGCTGGTGCATTTTTTTACAAACAGTTGGCACAGAGTATTTGTTGCTAATTGTGTTTCGTTTTCTTGAA...

TR20186 | c1_g1_i23 | Alexis.tri |1410.....1420.....1430.....1440.....1450.....1460.....1470.....
 ACTGATGAATTCAGGTACACCGACAAGCATGTTATCTTCACATTCATTGTCCTTGAATGTTAATAGTTGGCCC
 TR20186 | c1_g1_i7 | Alexis.trin | ACTGATGAATTCAGGTACACCGACAAGCATGTTATCTTCACATTCATTGTCCTTGAATGTTAATAGTTGGCCC
 TR20186 | c1_g1_i8 | Alexis.trin | ACTGATGAATTCAGGTACACCGACAAGCATGTTATCTTCACATTCATTGTCCTTGAATGTTAATAGTTGGCCC
 TR20186 | c1_g1_i1 | Alexis.trin | ACTGATGAATTCAGGTACACCGACAAGCATGTTATCTTCACATTCATTGTCCTTGAATGTTAATAGTTGGCCC
 TR20186 | c1_g1_i2 | Alexis.trin | ACTGATGAATTCAGGTACACCGACAAGCATGTTATCTTCACATTCATTGTCCTTGAATGTTAATAGTTGGCCC
 TR20186 | c1_g1_i10 | Alexis.tri | ACTGATGAATTCAGGTACACCGACAAGCATGTTATCTTCACATTCATTGTCCTTGAATGTTAATAGTTGGCCC
 TR20186 | c1_g1_i25 | Alexis.tri | ACTGATGAATTCAGGTACACCGACAAGCATGTTATCTTCACATTCATTGTCCTTGAATGTTAATAGTTGGCCC
 TR20186 | c1_g1_i28 | Alexis.tri | ACTGATGAATTCAGGTACACCGACAAGCATGTTATCTTCACATTCATTGTCCTTGAATGTTAATAGTTGGCCC
 TR20186 | c1_g1_i6 | Alexis.trin | ACTGATGAATTCAGGTACACCGACAAGCATGTTATCTTCACATTCATTGTCCTTGAATGTTAATAGTTGGCCC
 TR20186 | c1_g1_i14 | Alexis.tri | ACTGATGAATTCAGGTACACCGACAAGCATGTTATCTTCACATTCATTGTCCTTGAATGTTAATAGTTGGCCC
 TR20186 | c1_g1_i13 | Alexis.tri | ACTGATGAATTCAGGTACACCGACAAGCATGTTATCTTCACATTCATTGTCCTTGAATGTTAATAGTTGGCCC
 TR20186 | c1_g1_i20 | Alexis.tri | ACTGATGAATTCAGGTACACCGACAAGCATGTTATCTTCACATTCATTGTCCTTGAATGTTAATAGTTGGCCC
 TR20186 | c1_g1_i3 | Alexis.trin | ACTGATGAATTCAGGTACACCGACAAGCATGTTATCTTCACATTCATTGTCCTTGAATGTTAATAGTTGGCCC
 TR20186 | c1_g1_i11 | Alexis.tri | ACTGATGAATTCAGGTACACCGACAAGCATGTTATCTTCACATTCATTGTCCTTGAATGTTAATAGTTGGCCC
 TR20186 | c1_g1_i22 | Alexis.tri | ACTGATGAATTCAGGTACACCGACAAGCATGTTATCTTCACATTCATTGTCCTTGAATGTTAATAGTTGGCCC
 TR20186 | c1_g1_i24 | Alexis.tri | ACTGATGAATTCAGGTACACCGACAAGCATGTTATCTTCACATTCATTGTCCTTGAATGTTAATAGTTGGCCC
 TR20186 | c1_g1_i18 | Alexis.tri | ACTGATGAATTCAGGTACACCGACAAGCATGTTATCTTCACATTCATTGTCCTTGAATGTTAATAGTTGGCCC
 TR10524 | c1_g1_i8 | Franka.trin |CATTTGTCCTTGAATGTTAATAGTTGGCCC
 TR24411 | c0_g2_i13 | Esterel.tr | ACTGATGAATTCAGGTACACCGGCAAGCATGTTATCTTCACATCCATTGTCCTTGAATGTTAATAGTTGGCCC
 TR24411 | c0_g2_i5 | Esterel.tri |
 TR24411 | c0_g2_i12 | Esterel.tr | ACTGATGAATTCAGGTACACCGGCAAGCATGTTATCTTCACATCCATTGTCCTTGAATGTTAATAGTTGGCCC
 TR10524 | c1_g1_i3 | Franka.trin |CATTTGTCCTTGAATGTTAATAGTTGGCCC
 TR24411 | c0_g2_i6 | Esterel.tri | ACTGATGAATTCAGGTACACCGACAAGCATGTTATCTTCACATTCATTGTCCTTGAATGTTAATAGTTGGCCC

Figure 10: Fragment of alignment produced by `annotate_cluster.pl`, rendered with [BioEdit](#).

Finally, option `-c` can be invoked to collapse aligned sequences from the same species or taxon. This might be useful when working with clusters of transcript isoforms, which are often redundant and broken in possibly overlapping fragments. Taking the same example cluster, we could try to collapse isoforms with overlaps ≥ 30 residues like this:

```
./annotate_cluster.pl -f outdir/1004_TR425-c0_g2_i1.fna -o 1004_TR425-c0_g2_i1.aln.fna -D -c 30
```

Note that by default the script does not tolerate mismatches between sequences to be collapsed, but that behaviour can be relaxed by editing the value of variable `$MAXMISMCOLLAP=0` at the top of the script. Instead, as BLASTN-placed gaps in identical sequences can often move, by default two such gaps are accepted (see variable `MAXGAPSCOLLAP=2`).

By default, the script `annotate_cluster.pl` looks for the longest sequences and aligns all other cluster sequences to it with BLASTN (megablast). The user can also pass an external, reference sequence to guide cluster alignment (see option `-r`). However, in either case, clusters of transcripts often contain a fraction of BLASTN hits that do not match the longest/reference sequence; instead, they align towards the 5' or 3' of other sequences of the clusters and are thus not included in the produced multiple sequence alignment (MSA):

```
-----<= longest/reference sequence
      -----
-----
-----
-----
      ....<= sequences not included in MSA
      ..
```

5 A step-by-step protocol with barley assembled transcripts

This section describes the steps required to proceed with the analysis of barley transcripts with folder `test_barley`, which you should get with the software. The following commands are to be pasted in your terminal:

```
## set get_homologues path if not already in $PATH
export GETHOMS=~/.soft/github/get_homologues/

cd test_barley

## 1) prepare sequences
cd seqs

# download all transcriptomes
wget -c -i wgetlist.txt

# extract CDS sequences (this takes several hours)
# choose cdsCPP.sh if dependency Inline::CPP is available in your system
# the script will use 20 CPU cores, please adapt it to your system
./cds.sh

# clean and compress
#rm -f *_* *noORF* *transcript*
#gzip *diamond*

# put cds sequences aside
mv *cds.f*gz ../cds
cd ..

# check lists of accessions are in place (see HOWTO.txt there)
ls cds/*list

## 2) cluster sequences and start the analyses

# calculate protein domain frequencies (Pfam)
$GETHOMS/get_homologues-est.pl -d cds -D -m cluster -o &> log.cds.pfam

# alternatively, if not running in a SGE cluster, taking for instance 20 CPUs
$GETHOMS/get_homologues-est.pl -d cds -D -n 20 -o &> log.cds.pfam

# calculate 'control' cds clusters
$GETHOMS/get_homologues-est.pl -d cds -M -t 0 -m cluster &> log.cds

# get non-cloud clusters
$GETHOMS/get_homologues-est.pl -d cds -M -t 3 -m cluster &> log.cds.t3

# clusters for dN/dS calculations
$GETHOMS/get_homologues-est.pl -d cds -e -M -t 4 -m cluster &> log.cds.t4.e

# leaf clusters and pangenome growth simulations with soft-core
$GETHOMS/get_homologues-est.pl -d cds -c -z \
  -I cds/leaf.list -M -t 3 -m cluster &> log.cds.leaf.t3.c

# produce pan-genome matrix and allocate clusters to occupancy classes

# all occupancies
```

```

$GETHOMS/compare_clusters.pl -d cds_est_homologues/Alexis_0taxa_algOMCL_e0_ \
-o clusters_cds -m -n &> log.compare_clusters.cds

# excluding cloud clusters, the most unreliable in our benchmarks
$GETHOMS/compare_clusters.pl -d cds_est_homologues/Alexis_3taxa_algOMCL_e0_ \
-o clusters_cds_t3 -m -n &> log.compare_clusters.cds.t3
$GETHOMS/parse_pangenome_matrix.pl -m clusters_cds_t3/pangenome_matrix_t0.tab -s \
&> log.parse_pangenome_matrix.cds.t3

# make pan-genome growth plots
$GETHOMS/plot_pancore_matrix.pl -i cds_est_homologues/core_genome_leaf.list_algOMCL.tab \
-f core_both &> log.core.plots
$GETHOMS/plot_pancore_matrix.pl -i cds_est_homologues/pan_genome_leaf.list_algOMCL.tab \
-f pan &> log.pan.plots

## 3) annotate accessory genes

# find [-t 3] SBCC073 clusters absent from references
$GETHOMS/parse_pangenome_matrix.pl -m clusters_cds_t3/pangenome_matrix_t0.tab \
-A cds/SBCC073.list -B cds/ref.list -g &> log.acc.SBCC073
mv clusters_cds_t3/pangenome_matrix_t0__pangenes_list.txt \
clusters_cds_t3/SBCC073_pangenes_list.txt

# how many SBCC073 clusters are there?
perl -lane 'if($F[0] =~ /SBCC073/){ foreach $c (1 .. $#F){ if($F[$c]>0){ $t++ } }; print $t }' \
clusters_cds_t3/pangenome_matrix_t0.tab

# find [-t 3] Scarlett clusters absent from references
$GETHOMS/parse_pangenome_matrix.pl -m clusters_cds_t3/pangenome_matrix_t0.tab \
-A cds/Scarlett.list -B cds/ref.list -g &> log.acc.Scarlett
mv clusters_cds_t3/pangenome_matrix_t0__pangenes_list.txt \
clusters_cds_t3/Scarlett_pangenes_list.txt

# find [-t 3] H.spontaneum clusters absent from references
$GETHOMS/parse_pangenome_matrix.pl -m clusters_cds_t3/pangenome_matrix_t0.tab \
-A cds/spontaneum.list -B cds/ref.list -g &> log.acc.spontaneum
mv clusters_cds_t3/pangenome_matrix_t0__pangenes_list.txt \
clusters_cds_t3/spontaneum_pangenes_list.txt

# Pfam enrichment tests

# core
$GETHOMS/pfam_enrich.pl -d cds_est_homologues -c clusters_cds -n \
-x clusters_cds_t3/pangenome_matrix_t0__core_list.txt -e -p 1 \
-r SBCC073 > SBCC073_core.pfam.enrich.tab

$GETHOMS/pfam_enrich.pl -d cds_est_homologues -c clusters_cds -n \
-x clusters_cds_t3/pangenome_matrix_t0__core_list.txt -e -p 1 \
-r SBCC073 -t less > SBCC073_core.pfam.deplet.tab

# accessory
$GETHOMS/pfam_enrich.pl -d cds_est_homologues -c clusters_cds -n \
-x clusters_cds_t3/SBCC073_pangenes_list.txt -e -p 1 -r SBCC073 \
-f SBCC073_accessory.fna > SBCC073_accessory.pfam.enrich.tab

```

```

$GETHOMS/pfam_enrich.pl -d cds_est_homologues -c clusters_cds -n \
-x clusters_cds_t3/Scarlett_pangenomes_list.txt -e -p 1 -r Scarlett \
-f Scarlett_accessory.fna > Scarlett_accessory.pfam.enrich.tab

$GETHOMS/pfam_enrich.pl -d cds_est_homologues -c clusters_cds -n \
-x clusters_cds_t3/spontaneum_pangenomes_list.txt -e -p 1 -r Hs_ \
-f spontaneum_accessory.fna > spontaneum_accessory.pfam.enrich.tab

# note that output files contain data such as the mean length of sequences

# get merged stats for figure
perl suppl_scripts/_add_Pfam_domains.pl > accessory_stats.tab
perl -lane 'print if($F[0] >= 5 || $F[1] >= 5 || $F[2] >= 5)' \
    accessory_stats.tab > accessory_stats_min5.tab
Rscript suppl_scripts/_plot_heatmap.R

```


6 Frequently asked questions (FAQs)

Please see also the FAQs in [manual_get_homologues.pdf](#).

- What's the performance gain of v2?

After evolving parts of the original code base, and fixing some bugs (see CHANGES.txt), both *get_homologues.pl* and *get_homologues-est.pl* have significantly improved their performance, as can be seen in the figure, which combines data from the original benchmark and new data generated after v2 was in place.




Figure 11: Computing time and RAM requirements of the original algorithm (OMCL, measured on 6 sequence sets) as compared to the updated v2 code (measured on 3 three sets).

- What are the main caveats when clustering transcripts/CDS sequences?

get_homologues-est.pl has been mainly tested with plant sequences, using both CDS sets from whole-genome annotations and also transcripts from expression experiments. The main problems we have found so far are split genes, frequent artifacts in genome assemblies, incomplete genes which lack exons, for the same previous reasons, and retained introns, which are common among plant transcripts. These three common situations are illustrated in the figure.

- What are those chimeras warnings produced when running *transcripts2cds*?




Figure 12: Common problems faced when clustering transcripts/CDS sequences.

When subroutine `transcripts::parse_blastx_cds_sequences` reads BLASTX/DIAMOND results checks whether secondary alignment starts

- Why have you not implemented the COG algorithm in the `get_homologues-est.pl`?

We have left the COG algorithm out of `get_homologues-est.pl` as it will take some more work to integrate it with redundant isoform calling, which is important for EST datasets. However, it should be possible to do it.

- The number of clusters produced with `-C 75 -S 85` does not match the pangenome/pantranscriptome size estimated with option `-c`

The reason for these discrepancies is that these are fundamentally different analyses. While the default runmode simply groups sequences trying to put in the same cluster isoforms of orthologues and very close inparalogues, a genome composition analysis performs a simulation in order to estimate how many novel sequences are added by genomes/transcriptomes sampled in random order. In terms of code, there are a couple of key global variables set in `lib/marfil_homology.pm`, lines 135-138, which control how a gene/transcript is compared to previously processed sequences in order to call it novel:

```
$MIN_PERSEQID_HOM_EST = 70.0;
$MIN_COVERAGE_HOM_EST = 50.0;
```

These values are equivalent to say that any sequence with *coverage* $\geq 50\%$ and *identity* $\geq 70\%$ to previous genes/transcripts will be considered simply a homologue and won't be accumulated to the growing pangenome/pantranscriptome. You might want to change these values to increase or relax the stringency and to match the parameters set to produce your clusters.

- Can `get_homologues-est.pl` be used to analyze non-coding sequences?

In principle the software should work with any type of nucleotide sequences. For instance, the next figure shows how it can be used to analyze conserved non-coding sequences among *Brachypodium distachyon* and rice, with a median BLASTN alignment length of 32.

- I produced 2 different parsimony trees with `compare_clusters.pl`, is there a way to merge them and add bootstrap values?




Figure 13: Core-genome composition analysis of conserved non-coding sequences (CNS) from 56 *Brachypodium distachyon* ecotypes and rice.

The two trees are equally parsimonious, that is, have the same parsimony score but different topologies. One possibility to combine them into one topology is to compute either the majority rule consensus (mjr) tree, for example with *consense* from the PHYLIP package, or represent a network consensus with a program such as *splitstree*.

Regarding the bootstrapping, you could write some R code (or in any other language) to randomly sample the columns in the pangenome matrix with replacement to construct a new, bootstrapped matrix with the same number of columns as the original one. You should generate 100 or 500 of these matrices and then run *pars* on each one of them. Then run *consense* to obtain the mjr consensus tree and associated split frequencies (bootstrap support values for each bipartition). An easy way of achieving this would be with R's *boot* package (see [example](#)).

Another option is to call *seqboot* from the PHYLIP package to generate N bootstrap pseudo-replicates of the matrix. Rename the resulting file as *infile* and call *pars* to read the file, using the option *m no_of_pseudoreplicates* to run the standard (Fitch) parsimony analysis on each of the bootstrapped matrices. PHYLIP *pars* will generate an outtree file containing as many trees as bootstrapped matrices found in *infile*. Rename *outtree* to *intree* and call *consense* to generate the default majority rule consensus tree. This tree is only a cladogram (only topology, no branch lengths). The node labels correspond to the number of bootstrap pseudoreplicates in which that particular bipartition was found.

- How can I produce a maximum likelihood (ML) tree with bootstrap values from a pangenome matrix?

Script `./compare_clusters.pl -m` produces three versions of the pangenome matrix: `.tab`, `.phy` and `.fasta`. This FASTA file can be analyzed with ML software such as [IQ-TREE](#) (PubMed=[25371430](#)) both online or in the terminal with a command such as:

```
path_to_iqtree -s pangenome_matrix_t0.fasta -st BIN -m TEST -bb 1000 -alrt 1000
```

This will produce an optimal ML tree after selecting a binary substitution model with both bootstrap and aLRT support values.

- Is there a way to plot ANI matrices of soft-core clusters?

Let's say you have 20 genomes, then 95% of them are exactly 19 taxa, which is the minimum occupancy that defines soft-core clusters (see global variable `$SOFTCOREFRACTION`). You should then compute the ANI matrix as follows:

```
./get_homologues.pl -d your_data -a 'CDS' -A -M -t 19
```

And then plot the resulting matrix with script *hcluster_matrix.sh*.

- When I use the *hcluster_matrix.sh* script the trees in the output of the newick file and the heatmap differ. Is there a reason for this?

The difference in the topologies of the NJ trees and the row-dendrogram of the heatmaps differ because the heatmaps are ordered bi-dimensionally. That is, the heatmap plot shows only the row-dendrogram, but the matrix is ordered also by columns. The NJ tree is computed from the distance matrix that you indicate the program to calculate for you (ward.D2 is the default).

- I ran *get_homologues* with fasta files of 78 genomes; is there a way to export a 78 x 78 matrix of the number of homologues shared between each genome?

If you did your analysis requesting cluster of all occupancies (`-t 0`) then you can get what you want in two steps. First, you must produce a pangenome matrix with `compare_clusters -d ... -m`. Now it is possible to request an intersection pangenome matrix (`pangenome_matrix_t0__intersection.tab`) which contains the number of sequence clusters shared by any two pairs of genomes with `parse_pangenome_matrix.pl -m pangenome_matrix_t0.tab -s -x`. Note that these clusters might contain several inparalogues of the same species.

7 Credits and references

get_homologues-est.pl is designed, created and maintained at the [Laboratory of Computational Biology](#) at Estación Experimental de Aula Dei/CSIC in Zaragoza (Spain) and at the [Center for Genomic Sciences](#) of Universidad Nacional Autónoma de México (CCG/UNAM).

The code was written mostly by Bruno Contreras-Moreira and Pablo Vinuesa, but it also includes code and binaries from [OrthoMCL v1.4](#) (algorithm OMCL, -M), [NCBI Blast+](#), [MVIEW](#), [DIAMOND](#) and [BioPerl 1.5.2](#).

Other contributors: Carlos P Cantalapiedra, Roland Wilhelm.

We ask the reader to cite the main reference describing the *get_homologues* software,

- Contreras-Moreira B, Cantalapiedra CP, Garcia Pereira MJ, Gordon S, Vogel JP, Igartua E, Casas AM and Vinuesa P (2017) Analysis of plant pan-genomes and transcriptomes with *GET_HOMOLOGUES-EST*, *a clustering solution for sequences of the same* and also the original papers describing the included algorithms and databases, accordingly:
 - Li L, Stoeckert CJ Jr, Roos DS (2003) OrthoMCL: identification of ortholog groups for eukaryotic genomes. *Genome Res.* 13(9):2178-89.
 - Altschul SF, Madden TL, Schaffer AA, Zhang J, Zhang Z, Miller W and Lipman DJ (1997) Gapped BLAST and PSI-BLAST: a new generation of protein database search programs. *Nucl. Acids Res.* 25(17): 3389-3402.
 - Stajich JE, Block D, Boulez K, Brenner SE, Chervitz SA, Dagdigian C, Fuellen G, Gilbert JG, Korf I, Lapp H, Lehvsilaiho H, Matsalla C, Mungall CJ, Osborne BI, Pocock MR, Schattner P, Senger M, Stein LD, Stupka E, Wilkinson MD, Birney E. (2002) The Bioperl toolkit: Perl modules for the life sciences. *Genome Res.* 12(10):1611-8.
 - hmmscan :: search sequence(s) against a profile database HMMER 3.1b2 (Feb 2015) <http://hmmer.org> Copyright (C) 2015 Howard Hughes Medical Institute. Freely distributed under the GNU General Public License (GPLv3).
 - Finn RD, Coghill P, Eberhardt RY, Eddy SR, Mistry J, Mitchell AL, Potter SC, Punta M, Qureshi M, Sangrador-Vegas A, Salazar GA, Tate J, Bateman A. (2016) The Pfam protein families database: towards a more sustainable future. *Nucleic Acids Res.* 44(D1):D279-85
 - Haas BJ, Papanicolaou A, Yassour M et al. (2013) De novo transcript sequence reconstruction from RNA-seq using the Trinity platform for reference generation and analysis. *Nat Protoc.* 8(8):1494-512.
 - Brown NP, Leroy C, Sander C (1998) MView: A Web compatible database search or multiple alignment viewer. *Bioinformatics.* 14 (4):380-381.
 - Buchfink B, Xie C, Huson DH (2015) Fast and sensitive protein alignment using DIAMOND. *Nat Methods.* 12(1):59-60

If you use the accompanying scripts the following references should also be cited:

- R Core Team (2013) R: A Language and Environment for Statistical Computing. <http://www.R-project.org> R Foundation for Statistical Computing, Vienna, Austria, ISBN3-900051-07-0